

Dataset used,  
S20\_Dataset\_Holidays.txt,

The screenshot shows a MobaXterm window with a terminal session on a remote host. The terminal displays the command `cat S20_Dataset_Holidays.txt` and its output, which is a list of 30 lines of data. Each line contains a country code, a city code, and a year, separated by commas. The data is as follows:

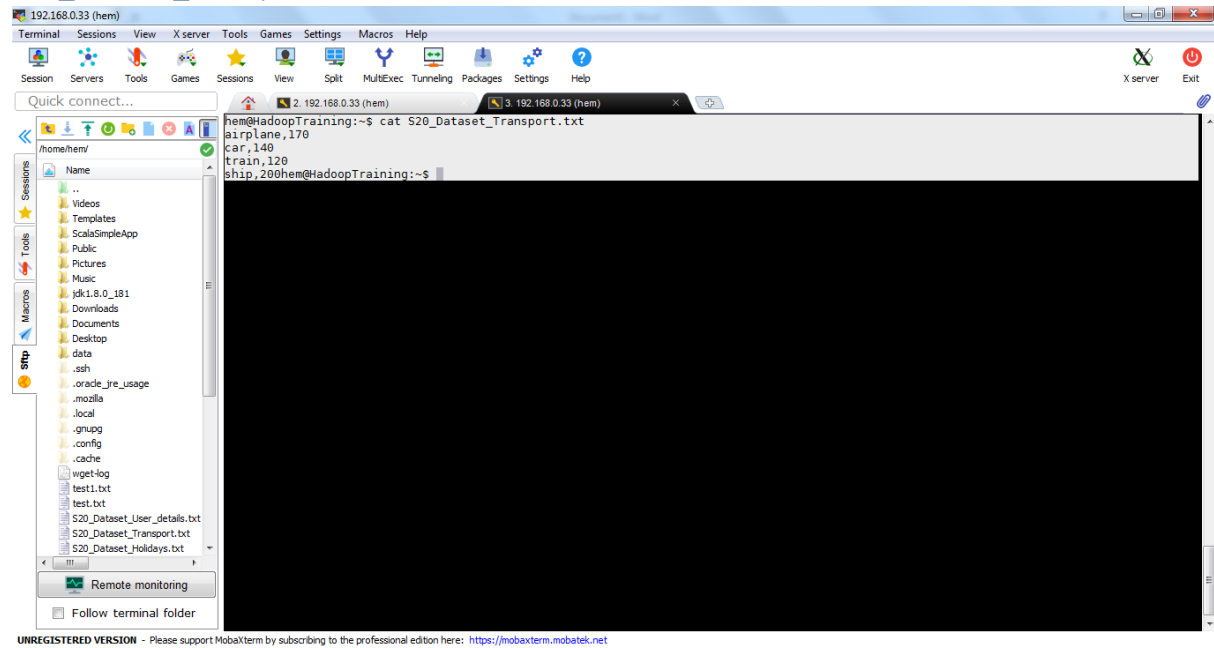
```
1,CHN,IND,airplane,200,1990
2,IND,CHN,airplane,200,1991
3,IND,CHN,airplane,200,1992
4,RUS,IND,airplane,200,1990
5,CHN,RUS,airplane,200,1992
6,AUS,PAK,airplane,200,1991
7,RUS,AUS,airplane,200,1990
8,IND,RUS,airplane,200,1991
9,CHN,RUS,airplane,200,1992
10,AUS,CHN,airplane,200,1993
1,AUS,CHN,airplane,200,1993
2,CHN,IND,airplane,200,1993
3,CHN,IND,airplane,200,1993
4,IND,AUS,airplane,200,1991
5,AUS,IND,airplane,200,1992
6,RUS,CHN,airplane,200,1993
7,CHN,RUS,airplane,200,1990
8,AUS,CHN,airplane,200,1990
9,IND,AUS,airplane,200,1991
10,RUS,CHN,airplane,200,1992
1,PAK,IND,airplane,200,1993
2,IND,RUS,airplane,200,1991
3,CHN,PAK,airplane,200,1991
4,CHN,PAK,airplane,200,1990
5,IND,PAK,airplane,200,1991
6,PAK,RUS,airplane,200,1991
7,CHN,IND,airplane,200,1990
8,RUS,IND,airplane,200,1992
9,RUS,IND,airplane,200,1992
10,CHN,AUS,airplane,200,1990
1,PAK,AUS,airplane,200,1993
5,CHN,PAK,airplane,200,1994
```

S20\_Dataset\_User\_details.txt,

The screenshot shows a MobaXterm window with a terminal session on a remote host. The terminal displays the command `cat S20_Dataset_User_details.txt` and its output, which is a list of 10 lines of data. Each line contains a name and an age, separated by a comma. The data is as follows:

```
1,mark,15
2,john,16
3,luke,17
4,lisa,27
5,mark,25
6,peter,22
7,james,21
8,andrew,55
9,thomas,46
10,annie,44
```

## S20\_Dataset\_Transport.txt



The screenshot shows a MobaXterm terminal window with a file explorer on the left. The terminal is open to a session named 'hem' at IP 192.168.0.33. The user has executed the command `cat S20_Dataset_Transport.txt`, and the output is displayed in the terminal window. The output shows a list of transportation modes and their counts: airplane,170; car,140; train,120; and ship,200. The file explorer on the left shows the directory structure of the user's home directory, including files like `S20_Dataset_Transport.txt`.

```
hem@HadoopTraining:~$ cat S20_Dataset_Transport.txt
airplane,170
car,140
train,120
ship,200hem@HadoopTraining:~$
```

## Task 1

1) What is the distribution of the total number of air-travelers per year

Spark Solution:-

```
val row = sc.textFile("/home/hem/S20_Dataset_Holidays.txt")

import org.apache.spark.storage.StorageLevel

row.persist(StorageLevel.MEMORY_ONLY)

val row1 = row.map(x => (x.split(",")(5).toInt,1))

val ans1 = row1.reduceByKey((x,y)=>(x+y)).foreach(println)
```

Output:-

```
scala> val row = sc.textFile("/home/hem/S20_Dataset_Holidays.txt")
row: org.apache.spark.rdd.RDD[String] = /home/hem/S20_Dataset_Holidays.txt MapPartitionsRDD[1] at textFile at <console>:24

scala> import org.apache.spark.storage.StorageLevel
import org.apache.spark.storage.StorageLevel

scala> row.persist(StorageLevel.MEMORY_ONLY)
res0: row.type = /home/hem/S20_Dataset_Holidays.txt MapPartitionsRDD[1] at textFile at <console>:24

scala> val row1 = row.map(x => (x.split(",")(5).toInt,1))
row1: org.apache.spark.rdd.RDD[(Int, Int)] = MapPartitionsRDD[2] at map at <console>:26

scala> val ans1 = row1.reduceByKey((x,y)=>(x+y)).foreach(println)
[Stage 0:>] (0 + 1) / 1(1994,1)
(1992,7)
(1990,8)
(1991,9)
(1993,7)
ans1: Unit = ()

scala>
```

2) What is the total air distance covered by each user per year

Spark Solution-:

```
val row = sc.textFile("/home/hem/S20_Dataset_Holidays.txt")
```

```
import org.apache.spark.storage.StorageLevel
```

```
row.persist(StorageLevel.MEMORY_ONLY)
```

```
val row1 = row.map(x => ((x.split(",")(0),x.split(",")(5)),x.split(",")(4).toInt))
```

```
val ans2 = row1.reduceByKey((x,y) => (x + y)).foreach(println)
```

Output-:

```
scala> val row = sc.textFile("/home/hem/S20_Dataset_Holidays.txt")
row: org.apache.spark.rdd.RDD[String] = /home/hem/S20_Dataset_Holidays.txt MapPartitionsRDD[5] at textFile at <console>:25

scala> import org.apache.spark.storage.StorageLevel
import org.apache.spark.storage.StorageLevel

scala> row.persist(StorageLevel.MEMORY_ONLY)
res1: row.type = /home/hem/S20_Dataset_Holidays.txt MapPartitionsRDD[5] at textFile at <console>:25

scala> val row1 = row.map(x => ((x.split(",")(0),x.split(",")(5)),x.split(",")(4).toInt))
row1: org.apache.spark.rdd.RDD[(String, String), Int] = MapPartitionsRDD[6] at map at <console>:28

scala> val ans2 = row1.reduceByKey((x,y) => (x + y)).foreach(println)
((3,1992),200)
((3,1993),200)
((5,1991),200)
((6,1991),400)
((10,1993),200)
((5,1992),400)
((8,1991),200)
((8,1990),200)
((1,1993),600)
((5,1994),200)
((2,1993),200)
((2,1991),400)
((4,1990),400)
((10,1992),200)
((3,1991),200)
((1,1990),200)
((10,1990),200)
((6,1993),200)
((9,1992),400)
((8,1992),200)
((7,1990),600)
((9,1991),200)
((4,1991),200)
ans2: Unit = ()

scala>
```

### 3) Which user has travelled the largest distance till date

Spark Solution-:

```
val row = sc.textFile("/home/hem/S20_Dataset_Holidays.txt")
```

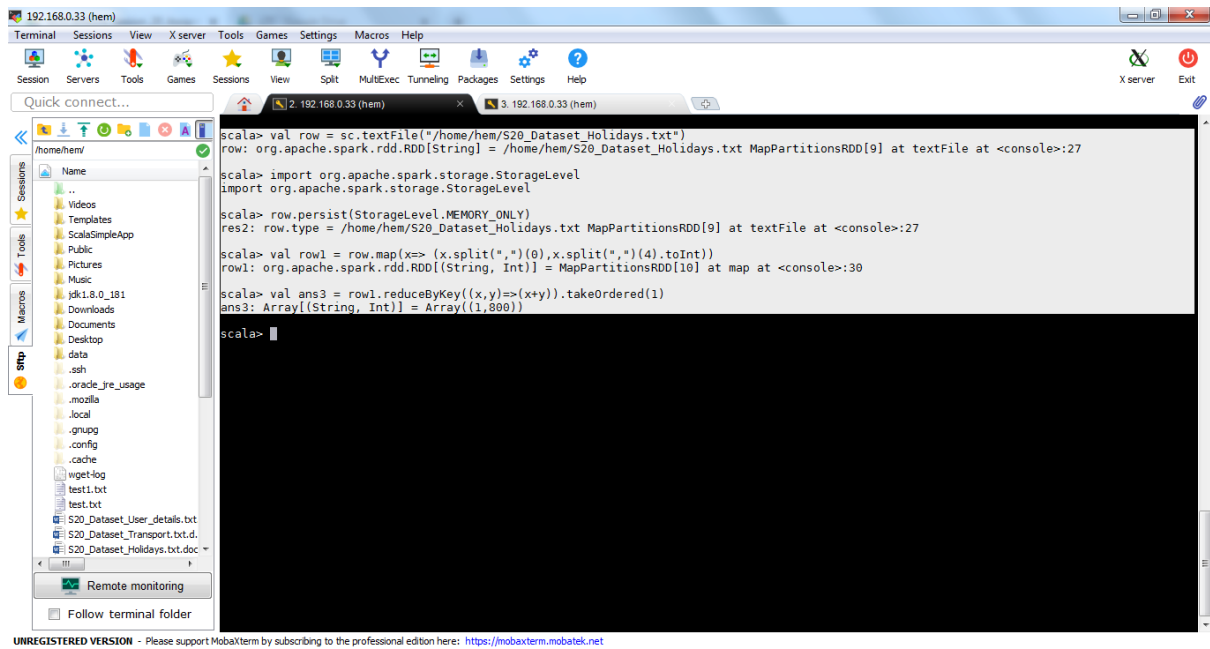
```
import org.apache.spark.storage.StorageLevel
```

```
row.persist(StorageLevel.MEMORY_ONLY)
```

```
val row1 = row.map(x=> (x.split(",")(0),x.split(",")(4).toInt))
```

```
val ans3 = row1.reduceByKey((x,y)=>(x+y)).takeOrdered(1)
```

Output-:



```
scala> val row = sc.textFile("/home/hem/S20_Dataset_Holidays.txt")
row: org.apache.spark.rdd.RDD[String] = /home/hem/S20_Dataset_Holidays.txt MapPartitionsRDD[9] at textFile at <console>:27

scala> import org.apache.spark.storage.StorageLevel
import org.apache.spark.storage.StorageLevel

scala> row.persist(StorageLevel.MEMORY_ONLY)
res2: row.type = /home/hem/S20_Dataset_Holidays.txt MapPartitionsRDD[9] at textFile at <console>:27

scala> val row1 = row.map(x=> (x.split(",")(0),x.split(",")(4).toInt))
row1: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[10] at map at <console>:30

scala> val ans3 = row1.reduceByKey((x,y)=>(x+y)).takeOrdered(1)
ans3: Array[(String, Int)] = Array((1,800))

scala>
```

4) What is the most preferred destination for all users.

Spark Solution:-

```
val row = sc.textFile("/home/hem/S20_Dataset_Holidays.txt")
```

```
import org.apache.spark.storage.StorageLevel
```

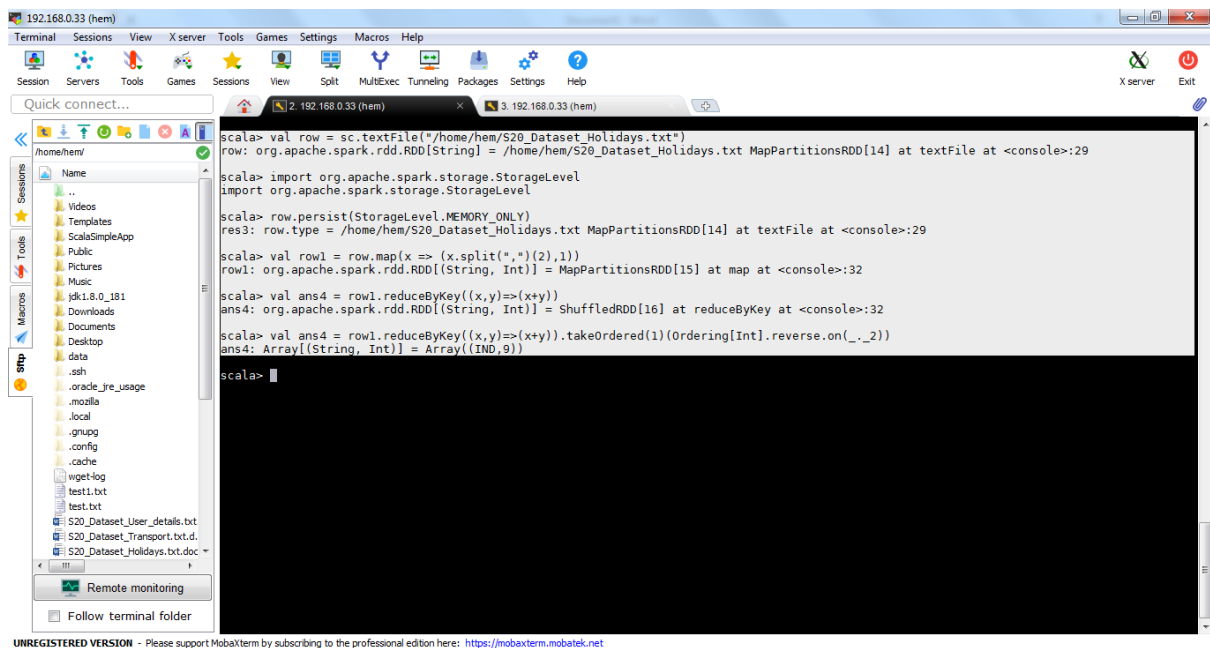
```
row.persist(StorageLevel.MEMORY_ONLY)
```

```
val row1 = row.map(x => (x.split(",")(2),1))
```

```
val ans4 = row1.reduceByKey((x,y)=>(x+y))
```

```
val ans4 = row1.reduceByKey((x,y)=>(x+y)).takeOrdered(1)(Ordering[Int].reverse.on(_._2))
```

Output:-



```
scala> val row = sc.textFile("/home/hem/S20_Dataset_Holidays.txt")
row: org.apache.spark.rdd.RDD[String] = /home/hem/S20_Dataset_Holidays.txt MapPartitionsRDD[14] at textFile at <console>:29

scala> import org.apache.spark.storage.StorageLevel
import org.apache.spark.storage.StorageLevel

scala> row.persist(StorageLevel.MEMORY_ONLY)
res3: row.type = /home/hem/S20_Dataset_Holidays.txt MapPartitionsRDD[14] at textFile at <console>:29

scala> val row1 = row.map(x => (x.split(",")(2),1))
row1: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[15] at map at <console>:32

scala> val ans4 = row1.reduceByKey((x,y)=>(x+y))
ans4: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[16] at reduceByKey at <console>:32

scala> val ans4 = row1.reduceByKey((x,y)=>(x+y)).takeOrdered(1)(Ordering[Int].reverse.on(_._2))
ans4: Array[(String, Int)] = Array((IND,9))

scala>
```

5) Which route is generating the most revenue per year

Spark Solution:-

```
val row1 = sc.textFile("/home/hem/hadoop/S20_Dataset_Holidays.txt")
```

```
val row2 = sc.textFile("/home/hem/hadoop/S20_Dataset_Transport.txt")
```

```
val row3 = sc.textFile("/home/hem/hadoop/S20
```

```
import org.apache.spark.storage.StorageLevel
```

```
row1.persist(StorageLevel.MEMORY_ONLY)_Dataset_User_details.txt")
```

```
row2.persist(StorageLevel.MEMORY_ONLY)
```

```
row3.persist(StorageLevel.MEMORY_ONLY)
```

Output:-

## 6) What is the total amount spent by every user on air-travel per year

Spark Solution-:

```
val row1 = sc.textFile("/home/hem/S20_Dataset_Holidays.txt")
val row2 = sc.textFile("/home/hem/S20_Dataset_Transport.txt")
val row3 = sc.textFile("/home/hem/S20_Dataset_User_details.txt")
import org.apache.spark.storage.StorageLevel
row1.persist(StorageLevel.MEMORY_ONLY)
row2.persist(StorageLevel.MEMORY_ONLY)
row3.persist(StorageLevel.MEMORY_ONLY)
val rowholiday =
row1.map(x=>(x.split(",")(0).toInt,x.split(",")(1),x.split(",")(2),x.split(",")(3),x.split(",")(4).toInt,x.split(",")(5).toInt))
val rowtransport = row2.map(x=>(x.split(",")(0),x.split(",")(1).toInt))
val rowuser = row3.map(x=>(x.split(",")(0).toInt,x.split(",")(1),x.split(",")(2).toInt))
val rowholidaysmap = rowholiday.map(x=>x._4->(x._2,x._5,x._6))
val rowtransportmap = rowtransport.map(x=>x._1->x._2)
val rowjoin1 = rowholidaysmap.join(rowtransportmap)
val rowroute = rowjoin1.map(x=>(x._2._1._1->x._2._1._3)->(x._2._1._2*x._2._2))
val rowrevenue = rowroute.groupByKey().map(x=>x._2.sum->x._1)
val ans = rowrevenue.sortByKey(false).first()
```

Output-:

```

scala> row1.persist(StorageLevel.MEMORY_ONLY)
res8: row1.type = /home/hem/S20_Dataset_Holidays.txt MapPartitionsRDD[46] at textFile at <console>:35

scala> row2.persist(StorageLevel.MEMORY_ONLY)
res9: row2.type = /home/hem/S20_Dataset_Transport.txt MapPartitionsRDD[48] at textFile at <console>:35

scala> row3.persist(StorageLevel.MEMORY_ONLY)
res10: row3.type = /home/hem/S20_Dataset_User_details.txt MapPartitionsRDD[50] at textFile at <console>:35

scala> val rowholiday = row1.map(x=>(x.split(",")(0).toInt,x.split(",")(1),x.split(",")(2),x.split(",")(3),x.split(",")(4).toInt,x.split(",")(5).toInt))
rowholiday: org.apache.spark.rdd.RDD[(Int, String, String, String, Int, Int)] = MapPartitionsRDD[51] at map at <console>:38

scala> val rowtransport = row2.map(x=>(x.split(",")(0),x.split(",")(1).toInt))
rowtransport: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[52] at map at <console>:38

scala> val rowuser = row3.map(x=>(x.split(",")(0).toInt,x.split(",")(1),x.split(",")(2).toInt))
rowuser: org.apache.spark.rdd.RDD[(Int, String, Int)] = MapPartitionsRDD[53] at map at <console>:38

scala> val rowholidaymap = rowholiday.map(x=>x._4->(x._2,x._5,x._6))
rowholidaymap: org.apache.spark.rdd.RDD[(String, (String, Int, Int))] = MapPartitionsRDD[54] at map at <console>:38

scala> val rowtransportmap = rowtransport.map(x=>x._1->x._2)
rowtransportmap: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[55] at map at <console>:38

scala> val rowjoin1 = rowholidaymap.join(rowtransportmap)
rowjoin1: org.apache.spark.rdd.RDD[(String, ((String, Int, Int), Int))] = MapPartitionsRDD[58] at join at <console>:40

scala> val rowroute = rowjoin1.map(x=>(x._2._1._1->x._2._1._3)->(x._2._1._2*x._2._2))
rowroute: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[59] at map at <console>:38

scala> val rowrevenue = rowroute.groupByKey().map(x=>x._2.sum->x._1)
rowrevenue: org.apache.spark.rdd.RDD[(Int, (String, Int))] = MapPartitionsRDD[61] at map at <console>:38

scala> val ans = rowrevenue.sortByKey(false).first()
ans: (Int, (String, Int)) = (204000,(IND,1991))

scala>

```

7) Considering age groups of < 20 , 20-35, 35 > ,Which age group is travelling the most every year.

Spark Solution:-

```

val row1 = sc.textFile("/home/hem/S20_Dataset_Holidays.txt")

val row2 = sc.textFile("/home/hem/S20_Dataset_Transport.txt")

val row3 = sc.textFile("/home/hem/S20_Dataset_User_details.txt")

import org.apache.spark.storage.StorageLevel

row1.persist(StorageLevel.MEMORY_ONLY)

row2.persist(StorageLevel.MEMORY_ONLY)

row3.persist(StorageLevel.MEMORY_ONLY)

val rowuser = row3.map(x=>(x.split(",")(0).toInt,x.split(",")(1),x.split(",")(2).toInt))

val rowholiday =
row1.map(x=>(x.split(",")(0).toInt,x.split(",")(1),x.split(",")(2),x.split(",")(3),x.split(",")(4).toInt,x.split(",")(5).toInt))

val ifElseMap = rowuser.map(x=>x._1->| {

| if(x._3<20)

| "20"

| else if(x._3>35)

| "35"

```



```
| else "20-35"
```

```
| })
```

```
val rowID = rowholiday.map(x => x._1 -> 1)
```

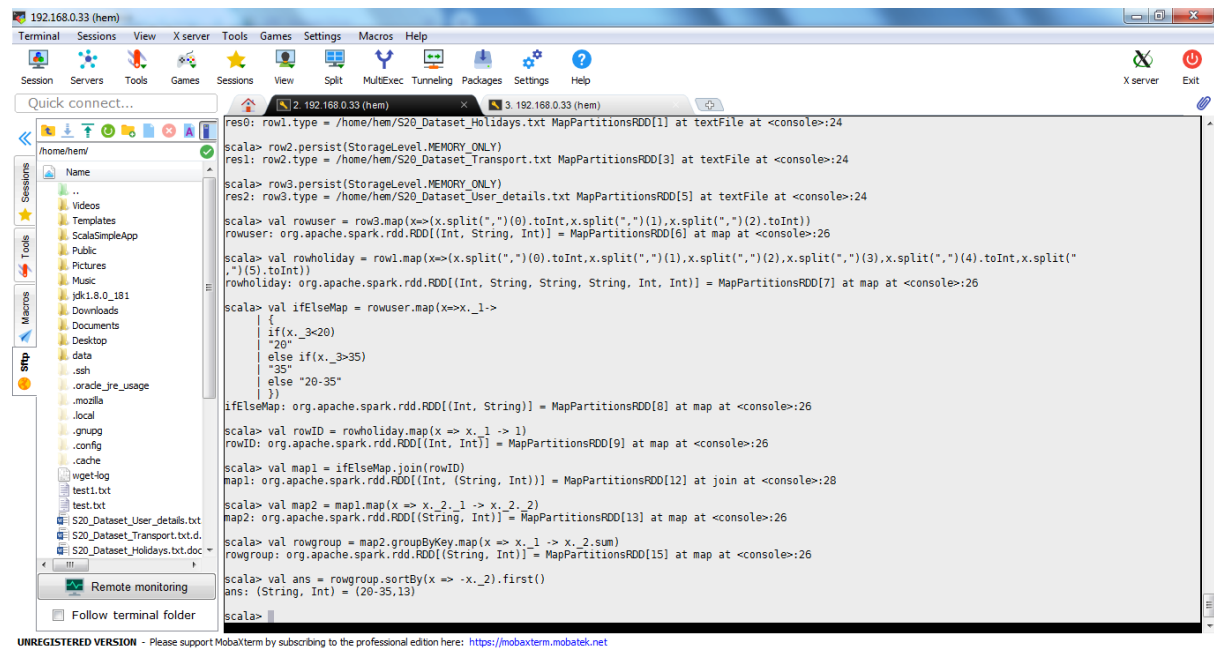
```
val map1 = ifElseMap.join(rowID)
```

```
val map2 = map1.map(x => x._2._1 -> x._2._2)
```

```
val rowgroup = map2.groupByKey.map(x => x._1 -> x._2.sum)
```

```
val ans = rowgroup.sortBy(x => -x._2).first()
```

Output-:



```
res0: row1.type = /home/hem/S20_Dataset_Holidays.txt MapPartitionsRDD[1] at textFile at <console>:24
scala> row2.persist(StorageLevel.MEMORY_ONLY)
res1: row2.type = /home/hem/S20_Dataset_Transport.txt MapPartitionsRDD[3] at textFile at <console>:24
scala> row3.persist(StorageLevel.MEMORY_ONLY)
res2: row3.type = /home/hem/S20_Dataset_User_details.txt MapPartitionsRDD[5] at textFile at <console>:24
scala> val rowuser = row3.map(x=>(x.split(",")(0).toInt,x.split(",")(1),x.split(",")(2).toInt))
rowuser: org.apache.spark.rdd.RDD[(Int, String, Int)] = MapPartitionsRDD[6] at map at <console>:26
scala> val rowholiday = row1.map(x=>(x.split(",")(0).toInt,x.split(",")(1),x.split(",")(2),x.split(",")(3),x.split(",")(4).toInt,x.split(",")(5).toInt))
rowholiday: org.apache.spark.rdd.RDD[(Int, String, String, String, Int, Int)] = MapPartitionsRDD[7] at map at <console>:26
scala> val ifElseMap = rowuser.map(x=>x._1->
| {
|   if(x._3<20)
|     "20"
|   else if(x._3>35)
|     "35"
|   else "20-35"
| })
ifElseMap: org.apache.spark.rdd.RDD[(Int, String)] = MapPartitionsRDD[8] at map at <console>:26
scala> val rowID = rowholiday.map(x => x._1 -> 1)
rowID: org.apache.spark.rdd.RDD[(Int, Int)] = MapPartitionsRDD[9] at map at <console>:26
scala> val map1 = ifElseMap.join(rowID)
map1: org.apache.spark.rdd.RDD[(Int, (String, Int))] = MapPartitionsRDD[12] at join at <console>:28
scala> val map2 = map1.map(x => x._2._1 -> x._2._2)
map2: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[13] at map at <console>:26
scala> val rowgroup = map2.groupByKey.map(x => x._1 -> x._2.sum)
rowgroup: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[15] at map at <console>:26
scala> val ans = rowgroup.sortBy(x => -x._2).first()
ans: (String, Int) = (20-35,13)
scala>
```

UNREGISTERED VERSION - Please support MobaXterm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net>